

Analysis of RNA-seq Data

Michael J. Guertin

April 15, 2020

This vignette provides step-by-step procedure for analysis of RNA-seq data.

Contents

1 RNA isolation and library preparation	3
1.1 Before beginning analysis	3
1.2 RNA extraction	3
1.3 Poly-A selection and ribosomal RNA depletion	3
1.4 Molecular biology of library preps	3
1.5 Small (17-35 nucleotides) RNA preps	4
1.6 Additional resources	5
1.7 Designing experiments	5
2 Retrieving and processing raw RNA-seq reads	6
2.1 Installing XCode, wget, homebrew, and SRA toolkit	6
2.1.1 Installing XCode toolkit	6
2.1.2 Installing homebrew	6
2.1.3 Installing wget	6
2.1.4 Installing SRA toolkit	6
2.1.5 Updating your \$PATH to find software	6
2.2 Finding relevant data and experiment details using ENCODE	7
2.3 Acquiring data from GEO	7
2.3.1 Processing SRA files	7
3 Aligning with HISAT2	9
3.1 Install HISAT2, samtools, and bedtools	9
3.1.1 Install hisat2	9
3.1.2 Install samtools	9
3.1.3 Install bedtools	9
3.2 Acquiring and processing the genome and annotation files	10
3.2.1 Retrieve the hg38.fa genome file	10
3.2.2 Build a set of genome index files with HISAT2	10
3.3 Aquiring and building gene annotations from GENCODE	10
3.4 Aligning with hisat2	10
3.5 Removing PCR duplicates	11
3.6 Converting files to upload into UCSC	11
3.7 Visualize peaks and bedGraph files with the UCSC browser	12
4 Counting reads in genes	14
4.1 Installing HTSeq	14
4.2 Count reads read in gene annotations	14
5 Running jobs on a cluster with SLURM	15
5.1 BASH script	15
5.2 SLURM script	16
5.3 Job Submission	17

6 Using R for analyses	18
6.1 Installing R packages on Rivanna	18
6.2 Importing counts data into R	18
6.2.1 Defining a function in R	19
6.3 Convert the Ensembl Gene (ENSG) name into a gene symbol	20
6.4 Differential expression with DESeq2	20
6.5 Gene Set Enrichment Analysis	24
7 Unexpected results	27
7.1 Modeling batch effect of the replicates	27
7.2 Exploratory analyses	29
7.3 Final analyses	32

List of Figures

1 RNA-seq library preparation	4
2 UCSC genome browser	12
3 Visualize your data in a browser.	13
4 PCA of RNA-seq analysis.	22
5 Gene Set Enrichment Gene Collections	24
6 Gene Set Enrichment plot	26
7 PCA after batch effect modeling.	28
8 Differential Expression MA plot	30
9 MA plot and GSEA of final analysis	33

1 RNA isolation and library preparation

1.1 Before beginning analysis

When designing an RNA-seq experiment, the first step is to determine what type of molecular biology workflow to perform for the sample preparation and the library preparation. Likewise, one needs to understand the molecular biology that generated the data of publicly available data.

1.2 RNA extraction

There are many different methods to generate RNA-seq data. RNA extraction is routine and kits are often used to extract total RNA. Although a DNase treatment is optional for many kits, I recommend this step.

1.3 Poly-A selection and ribosomal RNA depletion

Over 80% of the RNA in a given cell is ribosomal RNA (rRNA). Most researchers view rRNA as uninformative for the purposes of RNA-seq and there are two main methods to enrich for mRNA: 1) enrichment of polyadenylated (poly-A) RNAs—typically referred to as mRNA library prep; and 2) rRNA depletion—typically referred to as total RNA library prep. Both methods involve DNA/RNA hybridization to enrich or deplete the RNA species of interest. For mRNA preps, poly-A RNA is hybridized to oligo-dT beads and non-hybridized RNA is washed away. The positively selected poly-A RNA is then eluted and used as the input. For total RNA preparations, the ribosomal RNAs are hybridized to ssDNA that is complementary to the rRNA sequence. This negatively selects rRNA and the supernatant is the input for the library preparation. Alternatively this DNA/RNA hybrid is digested with RNaseH. The H refers to the RNase specificity—only digesting RNA in a DNA Hybrid.

1.4 Molecular biology of library preps

The molecular biology of library preps diverge depending upon whether the experimenter wishes to preserve the RNA strand information in the sequencing data. Unstranded libraries were common over a decade ago, but I see no reason why contemporary research would employ strand-ambiguous RNA-seq. Therefore, I will briefly discuss the workflow for the library preparation that I prefer and how clever molecular biologists developed methods to retain the strand specificity of the RNA throughout the library prep.

- 1) Isolate RNA.
- 2) Negative selection of rRNA.
- 3) Fragment the RNA with heat.
- 4) Anneal random hexamer DNA primers to the RNA.
- 5) Perform reverse transcription of the RNA using a reverse transcriptase that is engineered to operate at high temperatures. RT at higher temperatures relieves secondary structure of the RNA and leads to less biased libraries. Including actinomycin D at this step prevents second strand synthesis during the RT by inhibiting DNA-dependent DNA synthesis.
- 6) Perform second strand synthesis using the dUTP nucleotide as opposed to dTTP. dUTP incorporation is important for conferring strand specificity.
- 7) The next step is to blunt the ends of the DNA and generate a 5' A overhang. T4 Polymerase harbors 3' to 5' exonuclease activity and 5' to 3' polymerase activity, which will blunt both overhangs. T4 Polymerase also phosphorylates 5' DNA ends.
- 8) Taq polymerase is used to add an A to the 3' end of the dsDNA. The high throughput sequencing adapter is usually a forked Y adapter that contains a T overhang on the 5' end. This pairs with the A overhang, traditionally referred to as TA cloning. Alternatively, a hairpin adapter is used that incorporates a dUTP that effectively generates a Y adapter when the dUTP is digested.
- 9) Ligate the sequencing adapters with DNA ligase.
- 10) Digest dUTPs incorporated during the second strand synthesis with uracil-DNA glycosylase (UG-Dase). This confers sequence specificity because now you have a single strand DNA molecule with a distinct adapter on each end. The ssDNA represents the template strand.
- 11) The last step is to PCR amplify with a variety of multiplexed adapters that permit pooling of samples.

The dsDNA molecules can be sequenced using a specific primer for each end of the molecule. Refer to step 10 to determine which sequencing primer would result in the coding strand and which would be the reverse complement thereof.

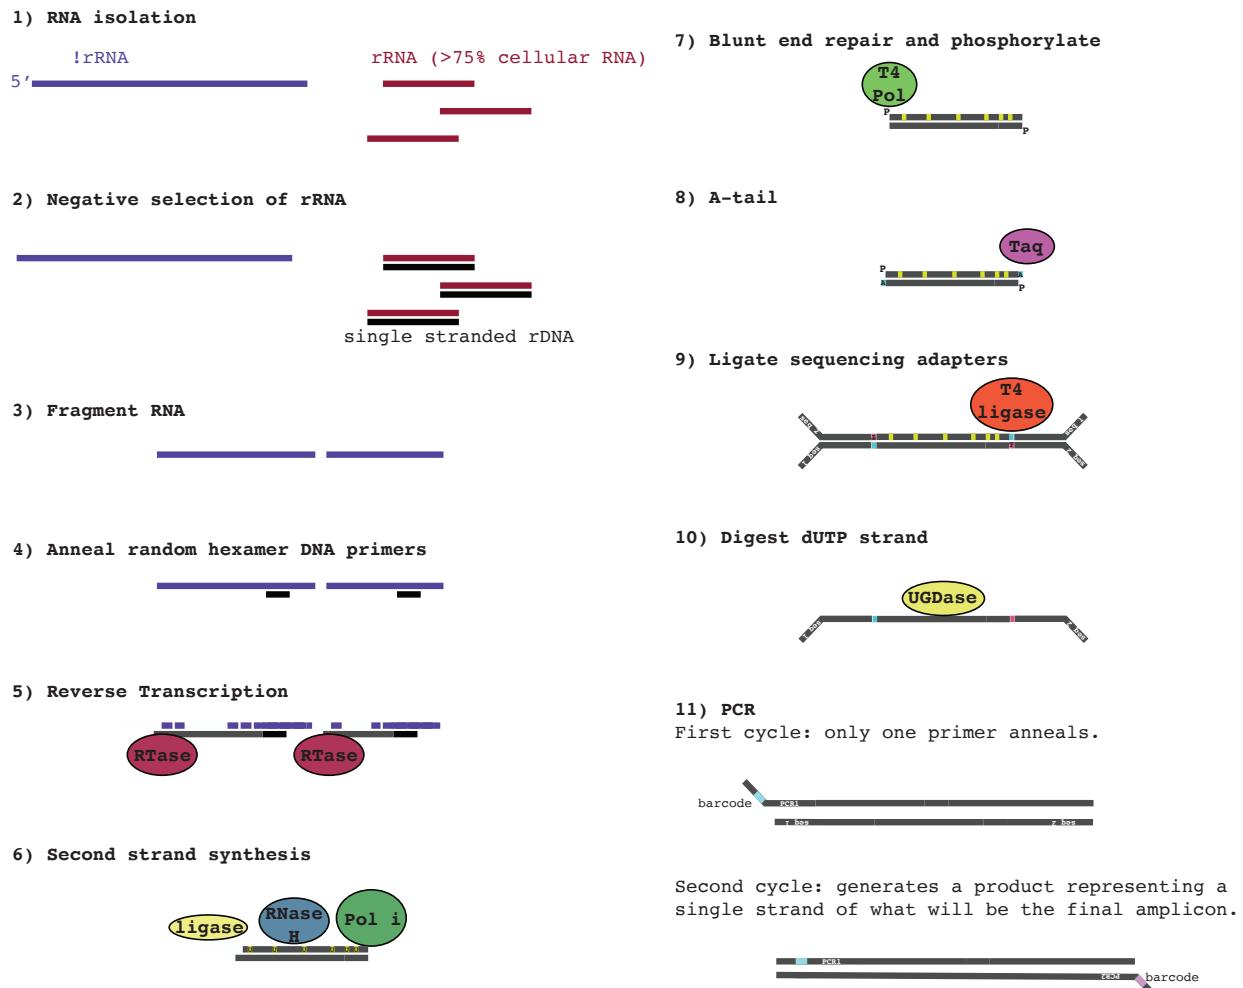


Figure 1: RNA-seq library preparation. 1) RNA isolation. 2) Negative selection of rRNA by hybridization with complementary rDNA and proceed with supernant or treat with RNase H. 3) Fragment RNA. 4) Anneal random hexamer DNA primers. 5) Perform reverse transcription. Note that Reverse Transcriptase has RNaseH activity, which partially digests the RNA template. 6) Second strand synthesis: *E. coli* DNA polymerase I is used for second strand synthesis. Partially digested RNA is used as the primers, RNase H degrades the RNA primers, and DNA polymerase fills the vacated primer sequence, and DNA ligase fills in the nicks. dUTP (yellow nucleotides) is included as opposed to dTTP. 7) T4 Polymerase generates blunt ends and phosphorylates 5' DNA ends. 8) Taq Polymerase A-tails the 3' DNA ends. 9) Y-primers are ligated with T4 ligase using TA cloning. 10) Uracil-DNA glycosylase digests the second strand. 11) PCR with barcoded adapters generates the final library.

1.5 Small (17-35 nucleotides) RNA preps

Even small RNA-seq libraries are generated from total RNA preparations, but a poly-A enrichment or rRNA depletion step is not necessary, as the libraries are size selected. For small RNA-seq libraries, one typically ligates the 3' adapter and the 5' adapter directly to RNAs isolated from cells. RNA isolation preserves the 3' Hydroxyl and 5' Phosphorylation modifications, so the input does not require enzymatic pre-treatment for ligation competency. The 3' adapter and the 5' adapter have distinct sequences, so strand specificity is preserved throughout the work flow.

1.6 Additional resources

The following is a good reference for the molecular biology behind different RNA-seq protocols:
https://www.neb.com/~media/nebus/files/brochures/nebnext_poster.pdf

1.7 Designing experiments

The second step, particularly when designing RNA-seq experiments, is to be aware of the features and limitations of your experimental design. A single experiment will provide the relative expression levels of RNAs in a particular cell type. This may be useful if you are looking to classify an unknown cell type by comparing to a panel of known transcription profiles of various tissues or cell types. This may be useful if your purpose is to annotate the transcriptome of a poorly characterized tissue/organism. Shotgun RNA-seq could also be used to characterize the genome of a novel RNA virus in order to develop a diagnostic test. However, it would be a challenge to use these data to develop new hypotheses to test. Another common implementation of RNA-seq is to compare expression between two conditions. This type of design will let you know which genes and isoforms are differentially expressed. These differentially expressed genes can be input to software that performs gene set enrichment analysis (GSEA) (Subramanian *et al.*, 2005) to identify biological pathways that are changing or to identify which transcription factor is likely regulating the differentially expressed genes (Wang *et al.*, 2018). One can begin to develop testable hypotheses from these data.

One can also use RNA-seq to test hypotheses, as outlined by this hypothetical example: One treats human cells with the plant hormone auxin and performs RNA-seq to determine which genes change expression. GSEA (Subramanian *et al.*, 2005) or Gene Ontology (Mi *et al.*, 2019) analysis indicates that canonical aryl hydrocarbon responsive genes are activated. BART (Wang *et al.*, 2018) identifies the Aryl Hydrocarbon Receptor (AHR) as the likely factor that regulates the differentially expressed genes. One can hypothesize that auxin activates AHR to regulate a set of genes. Test this hypothesis by depleting AHR, treating with auxin, and performing RNA-seq under the four conditions: -/+ AHR depletion and -/+ auxin treatment.

I use this as an example because we were using the auxin-inducible degron (AID) system and performed this control. We developed the hypothesis that auxin is activating the AHR and that this is a previously unappreciated deficiency of the AID system (Sathyan *et al.*, 2019). This brings us an exciting feature of RNA-seq: **we can develop hypotheses that were inconceivable prior to analyzing the data!**

2 Retrieving and processing raw RNA-seq reads

2.1 Installing XCode, wget, homebrew, and SRA toolkit

2.1.1 Installing XCode toolkit

Some of you will have personal Mac computers that require additional software installation. These code chunks will be contained within light blue boxes.

All Mac users should first install Xcode from the Mac App Store.

2.1.2 Installing homebrew

Homebrew is a convenient package manager and installer for macOS. Common Linux packages and some bioinformatics tools can be easily installed with homebrew. The website provides this installation code: <https://docs.brew.sh/Installation>. I do not recommend copying and pasting code from a PDF, because formatting and characters can be erroneously interpreted depending on your PDF reader/viewer. For instance, I cannot copy most code from Preview, but Adobe Acrobat works reasonably well. Note that generally these code chunks should be run successively, because the output of one chunk is often the input for a subsequent chunk.

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install.sh)"
```

2.1.3 Installing wget

The wget software package is used to retrieve files from the internet.

```
brew install wget
```

2.1.4 Installing SRA toolkit

fasterq-dump is an SRA tool (Kodama *et al.*, 2011) that extracts data directly from the Sequence Read Archive, where most high throughput sequencing data is deposited.

Software installation:

sra-tools: <https://github.com/ncbi/sra-tools>

Downloading the pre-compiled binaries that are compatible with your computer architecture is easiest (<https://github.com/ncbi/sra-tools/wiki/01.-Downloading-SRA-Toolkit>). Below we download the binaries for a Mac, move them to a directory that is present on most Mac computers, then update the computer's \$PATH so that Terminal will look in the proper place for the software.

```
wget http://ftp-trace.ncbi.nlm.nih.gov/sra/sdk/2.9.6-1/sratooldkit.2.9.6-1-mac64.tar.gz
tar -xzf sratooldkit.2.9.6-1-mac64.tar.gz
mv sratooldkit.2.9.6-1-mac64 /usr/local/bin/
export PATH="$PATH:/usr/local/bin/sratooldkit.2.9.6-1-mac64/bin"
```

2.1.5 Updating your \$PATH to find software

I add the PATH directly to my .bash_profile so that every time I open a new login shell the \$PATH variable updates. Alternatively, you could modify your .bashrc file, which runs on every interactive shell launch.

```
emacs ~/.bash_profile
```

Add the following line to the .bash_profile file:

```
export PATH="$PATH:/usr/local/bin/sratoolkit.2.9.6-1-mac64/bin"
```

Save and close emacs using the following keystrokes in sequence while keeping the previous keys engaged with each subsequent key (note that the "+" signs do not represent a key that should be pressed):
control + x + s
control + x + c

2.2 Finding relevant data and experiment details using ENCODE

The ENCODE consortium generated thousands of molecular genomics data sets that are publicly available (<https://www.encodecc.org>). The data below are from A549, which is a adenocarcinomic human alveolar basal epithelial cell line. The cells were treated with dexamethasone for various time points. All the available experimental information can be found on the ENCODE website:
<https://www.encodeproject.org/treatment-time-series/ENCSR897XFT/>. Below, I download the *mRNA-seq on A549 cell line treated with 100 nM dexamethasone for 0 hours AND 0.5 hours* for comparison. I encourage you to follow my workflow for the 0 minutes vs. 30 minutes comparison. For you adventurous types, I encourage you to choose an additional time point to analyze and compare against time zero.

2.3 Acquiring data from GEO

Sequencing files will come off the machine in FASTQ format, but they are stored as SRA (short read archive) files in data bases such as GEO (Gene Expression Omnibus) (<http://www.ncbi.nlm.nih.gov/geo/>).

All the files should be named with the cell type, conditions, description, the replicate number. No spaces. This makes it so that downstream analyses can be automated.

2.3.1 Processing SRA files

Retrieve and convert the binary SRA files to FASTQ files with the *sra-tools* kit; *fasterq-dump* automatically detects whether the data is single end or paired end. Paired end data will result in two output files and these data are paired end.

Back in Section 2.1.2 I recommended that you do not copy and paste from a PDF, but you may want the raw working shell and R scripts, so I uploaded them to GitHub and I will include the links. The following two code chunks are included in a single bash script: https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/section2.3.1.sh. This bash script can be accessed directly with *wget* and all URLs that begin with *raw.githubusercontent.com* are the raw code.

```
#no dex rep 1 (as defined by GEO description)
fasterq-dump SRR5093037
#no dex rep2
fasterq-dump SRR5093036
#no dex rep3
fasterq-dump SRR5093038
#no dex rep4
fasterq-dump SRR5093039

#30min 100nM dex rep 1
fasterq-dump SRR5093268
#30min 100nM dex rep 2
fasterq-dump SRR5093265
```

```
#30min 100nM dex rep 3  
fasterq-dump SRR5093267  
#30min 100nM dex rep 4  
fasterq-dump SRR5093266
```

Use the `mv` command to name the files something descriptive: cell line, treatment, replicate number, and paired end status.

```
mv SRR5093037_1.fastq A549_no_treatment_rep1_PE1.fastq  
mv SRR5093036_1.fastq A549_no_treatment_rep2_PE1.fastq  
mv SRR5093038_1.fastq A549_no_treatment_rep3_PE1.fastq  
mv SRR5093039_1.fastq A549_no_treatment_rep4_PE1.fastq  
  
mv SRR5093268_1.fastq A549_30min100nMdex_rep1_PE1.fastq  
mv SRR5093265_1.fastq A549_30min100nMdex_rep2_PE1.fastq  
mv SRR5093267_1.fastq A549_30min100nMdex_rep3_PE1.fastq  
mv SRR5093266_1.fastq A549_30min100nMdex_rep4_PE1.fastq  
  
mv SRR5093037_2.fastq A549_no_treatment_rep1_PE2.fastq  
mv SRR5093036_2.fastq A549_no_treatment_rep2_PE2.fastq  
mv SRR5093038_2.fastq A549_no_treatment_rep3_PE2.fastq  
mv SRR5093039_2.fastq A549_no_treatment_rep4_PE2.fastq  
  
mv SRR5093268_2.fastq A549_30min100nMdex_rep1_PE2.fastq  
mv SRR5093265_2.fastq A549_30min100nMdex_rep2_PE2.fastq  
mv SRR5093267_2.fastq A549_30min100nMdex_rep3_PE2.fastq  
mv SRR5093266_2.fastq A549_30min100nMdex_rep4_PE2.fastq  
  
gzip *fastq
```

The `gzcat` command uncompresses the gzipped file and writes the contents to stdout, which is piped to either `head` or `wc -l` in order to view the first ten lines and count the number of lines.

Note that each sequencing read has four lines associated with it:
Line 1 begins with a '@' and is followed by a unique identifier
Line 2 is the raw nucleotide sequence.
Line 3 begins with a '+' and is optionally followed by the same sequence identifier.
Line 4 encodes the quality scores for the sequence in Line 2, and must contain the same number of symbols as letters in the sequence. This line is a measure of how confident the base is called correctly.

```
gzcat A549_no_treatment_rep1_PE1.fastq.gz | head  
gzcat A549_no_treatment_rep1_PE1.fastq.gz | wc -l
```

3 Aligning with HISAT2

3.1 Install HISAT2, samtools, and bedtools

Commonly used software have extensive documentation and installation instructions, so if you run into challenges installing software, please refer to these resources.

3.1.1 Install hisat2

HISAT2 is commonly used alignment software for RNA-seq data (Kim *et al.*, 2015).

Software installation:

```
hisat2: https://daehwankimlab.github.io/hisat2/
```

```
wget https://cloud.biohpc.swmed.edu/index.php/s/hisat2-220-OSX_x86_64/download
mv download hisat2-2.2.0.zip
unzip hisat2-2.2.0.zip
mv hisat2-2.2.0 /usr/local/bin

export PATH="$PATH:/usr/local/bin/hisat2-2.2.0"
```

Refer to Section 2.1.5 to add this directory to your `~/.bash_profile`.

3.1.2 Install samtools

Samtools is an extensive suite of tools that allow you to interface with high-throughput sequencing data files in many ways.

Software installation:

```
samtools: http://www.htslib.org/download/
```

```
wget https://github.com/samtools/samtools/releases/download/1.10/samtools-1.10.tar.bz2
tar -xf samtools-1.10.tar.bz2
cd samtools-1.10
mkdir /usr/local/bin/samtools-1.10
./configure --prefix=/usr/local/bin/samtools-1.10
make
make install

export PATH=/usr/local/bin/samtools-1.10/bin:$PATH
```

Refer to Section 2.1.5 to add this directory to your `~/.bash_profile`.

3.1.3 Install bedtools

Bedtools utilities provide efficient and well-documented methods to perform *arithmetic* of genomic coordinate files. Originally developed at UVA by Aaron Quinlan and Ira Hall, this is the go to set of tools that most genomicists use to operate on genomic intervals.

Software installations:

```
bedtools: https://bedtools.readthedocs.io/en/latest/
```

```
brew install bedtools
```

3.2 Acquiring and processing the genome and annotation files

3.2.1 Retrieve the *hg38.fa* genome file

The first task is to retrieve the relevant genome from UCSC (Karolchik *et al.*, 2014): (<http://hgdownload.soe.ucsc.edu/downloads.html>). We want the *hg38.fa.gz* file. This is a zipped fasta file of the entire human genome.

```
wget http://hgdownload.soe.ucsc.edu/goldenPath/hg38/bigZips/hg38.fa.gz
gunzip hg38.fa.gz
```

Look at the format of a FASTA (pronounced with a hard A) file. Most of the sequence is either A,C,T, or G—but toward the telomeres there are many N's, due to the fact that assembling repetitive DNA is non-trivial.

```
head hg38.fa
```

3.2.2 Build a set of genome index files with HISAT2

Next we build the genome index with *hisat2* (Kim *et al.*, 2015). This task is performed once per genome.

```
hisat2-build hg38.fa hg38
```

3.3 Aquiring and building gene annotations from GENCODE

We built the indexed genome for *hisat2*, but now we will need to acquire and index the gene annotation GTF file. We will use GENCODE (<https://www.gencodegenes.org/human/>).

```
wget ftp://ftp.ebi.ac.uk/pub/databases/gencode/Gencode_human/release_33/gencode.v33.annotation.gtf.gz
gunzip gencode.v33.annotation.gtf.gz
python3 /usr/local/bin/hisat2-2.2.0/hisat2_extract_splice_sites.py gencode.v33.annotation.gtf > gencode.v33.splice.txt
```

3.4 Aligning with hisat2

Aligning using multiple processors in parallel speeds up the alignment. How many processors does your Mac have? I have 4 processors and I will leave one processor free by specifying the -p option as 3.

```
sysctl hw.ncpu | awk '{print $2}'
## 4
```

To align all the files in the same manner we can use shell scripting to loop over each file. The files are names descriptively, so we can extract the descriptive text from the input file (the \$name variable below) and use this to label the output files. The input files include the indexed hg38 genome (-x), the splice sites files, and each paired end FASTQ file. These files are used in the alignment and piped (|) to *samtools* to generate a BAM file (*hisat2* output is SAM format), and lastly this file is piped to a sorting function. SAM is a human-readable alignment file and BAM files are the binary counterpart, which stores data more efficiently, but is not human-readable text. Note that the \ after *gencode.v33.splice.txt* (and in other contexts throughout these shell chunks) is used to indicate that the line of code is extended onto the next line in the chunk.

```

for i in *PE1.fastq.gz
do
    name=$(echo $i | awk -F"_PE." '{print $1}')
    echo $name
    hisat2 -p 3 -x hg38 --rna-strandness FR --known-splicesite-infile gencode.v33.splice.txt \
        -1 $i -2 ${name}_PE2.fastq.gz | samtools view -bS - | samtools sort -n - -o $name.sorted.bam
done

```

This loop hard codes the hg38 prefix for the indexed files, the splice site file, and the number of processors. However, these will be changed depending upon the experiment, so these can be initialized and passed to the loop for more general usage. In the chunk below you can change the *idx*, *spliceFile*, and *processors* names in the header and these are passed to the loop. However, the loop still assumes that the data is paired end and named with the suffixes *_PE1.fastq.gz* and *_PE2.fastq.gz*. This chunk is uploaded as a bash script: https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/section3.4.sh.

```

idx=hg38
spliceFile=gencode.v33.splice.txt
processors=3

for i in *PE1.fastq.gz
do
    name=$(echo $i | awk -F"_PE." '{print $1}')
    echo $name
    hisat2 -p ${processors} -x ${idx} --rna-strandness FR --known-splicesite-infile ${spliceFile} \
        -1 ${i} -2 ${name}_PE2.fastq.gz | samtools view -bS - | samtools sort -n - -o $name.sorted.bam
done

```

3.5 Removing PCR duplicates

We can identify reads that start and end at the exact same genomic coordinates as presumptive PCR duplicates. This employs `samtools` once again and you will need to refer to `fixmate`, `sort`, and `markdup` within the `samtools` user manual to investigate the function of the pipe below. How does this `sort` command functionally differ than Section 3.4?

This code chunk is provided at: https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section3.5.sh. Note that this is not the `raw.githubusercontent.com` file, but you should be able to compare with the previous links to find the raw file.

```

for i in *.sorted.bam
do
    name=$(echo $i | awk -F".sorted.bam" '{print $1}')
    echo $name
    samtools fixmate -m $i - | samtools sort - | samtools markdup -rs - $name.rmPCRdup.bam
done

```

3.6 Converting files to upload into UCSC

You can also upload RNA-seq data to view in UCSC. You will use a position-sorted BAM file. Then use the `bedtools` function `genomeCoverageBed`. The zipped `bedGraph` can be uploaded to UCSC. The `BAM` file is sorted by position so the input to `genomeCoverageBed` is properly formatted. I am fastidious about how I want my data visualized in a browser, so the code chunk is a bit complicated. This code uses both paired end reads to visualize the data and generates a plus and minus track that are colored red and blue. I did leave out some code that normalizes each track, because it will further complicate the loop. But note that these are not scaled by sequencing depth. I typically use the `size.factors` calculated by `DESeq2` (not yet covered) and the `-scale` option in `genomeCoverageBed` to scale these browser tracks.

Code: https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section3.6.sh

```

for i in *.rmPCRdup.bam
do
    name=$(echo $i | awk -F".rmPCRdup.bam" '{print $1}')
    echo $name

```

```

samtools view -b -f 0x40 $i | samtools sort - -o ${name}.PE1.sorted
samtools view -b -f 0x80 $i | samtools sort - -o ${name}.PE2.sorted
genomeCoverageBed -bg -strand - -split -ibam ${name}.PE1.sorted | sortBed -i - > ${name}.pe1.minus.bedGraph
genomeCoverageBed -bg -strand + -split -ibam ${name}.PE1.sorted | sortBed -i - > ${name}.pe1.plus.bedGraph
genomeCoverageBed -bg -strand - -split -ibam ${name}.PE2.sorted | sortBed -i - > ${name}.pe2.minus.bedGraph
genomeCoverageBed -bg -strand + -split -ibam ${name}.PE2.sorted | sortBed -i - > ${name}.pe2.plus.bedGraph
cat ${name}.pe1.plus.bedGraph ${name}.pe2.plus.bedGraph | sortBed -i - > ${name}.merged.plus.bedGraph
touch temp.txt
echo "track type=bedGraph name=${name}.plus color=255,0,0 visibility=2" >> temp.txt
cat temp.txt ${name}.merged.plus.bedGraph > ${name}.plus.bedGraph
rm temp.txt
cat ${name}.pe1.minus.bedGraph ${name}.pe2.minus.bedGraph | sortBed -i - > ${name}.merged.minus.bedGraph
touch temp.txt
echo "track type=bedGraph name=${name}.minus color=0,0,255 visibility=2" >> temp.txt
cat temp.txt ${name}.merged.minus.bedGraph > ${name}.minus.bedGraph
rm temp.txt
rm *.merged.*us.bedGraph
rm *.pe*.us.bedGraph
gzip ${name}.minus.bedGraph
gzip ${name}.plus.bedGraph
done

```

3.7 Visualize peaks and bedGraph files with the UCSC browser

I prefer the UCSC browser to visualize my data and for exporting publication-quality figures (<https://genome.ucsc.edu>) (Karolchik *et al.*, 2014). Make sure you have the correct assembly, we are using the latest: hg38: <https://genome.ucsc.edu/cgi-bin/hgTracks?db=hg38>.

After clicking *manage custom tracks* (Figure 2), use *add custom tracks* to upload the *.bedGraph.gz files. Upload each file and customize the settings. You probably want to register and save sessions if you have your own data. Then you can save sessions and you will not need to upload the data more than once. I previously uploaded the data and saved a session: https://genome.ucsc.edu/s/Mike%20Guertin/hg38_class_example

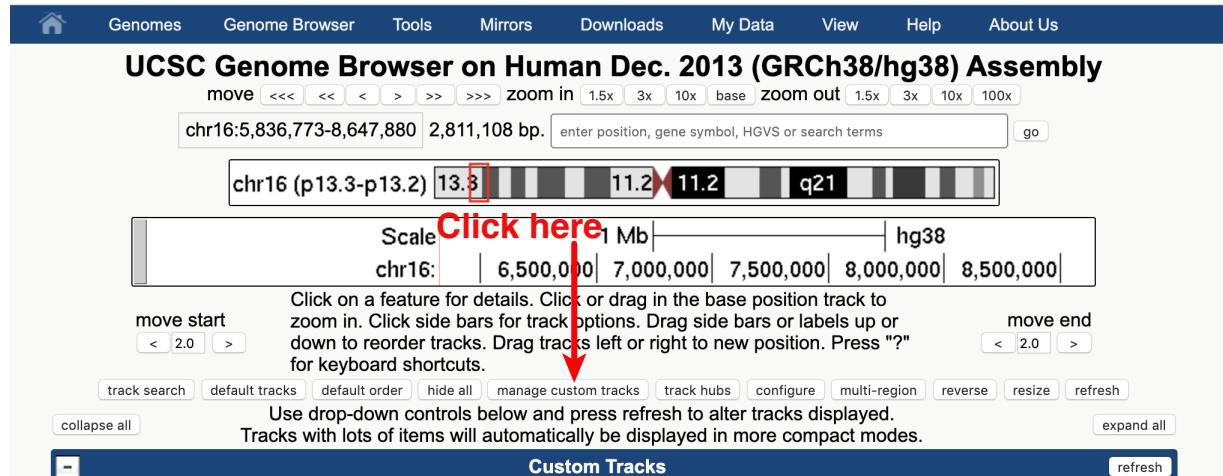


Figure 2: **UCSC genome browser.** There are many types of browsers one can use to visualize data, I like UCSC because it supports track hubs, sharing, and it does not run on the local computer.

Looking at your data in a browser can save you hours, days, or longer in troubleshooting.

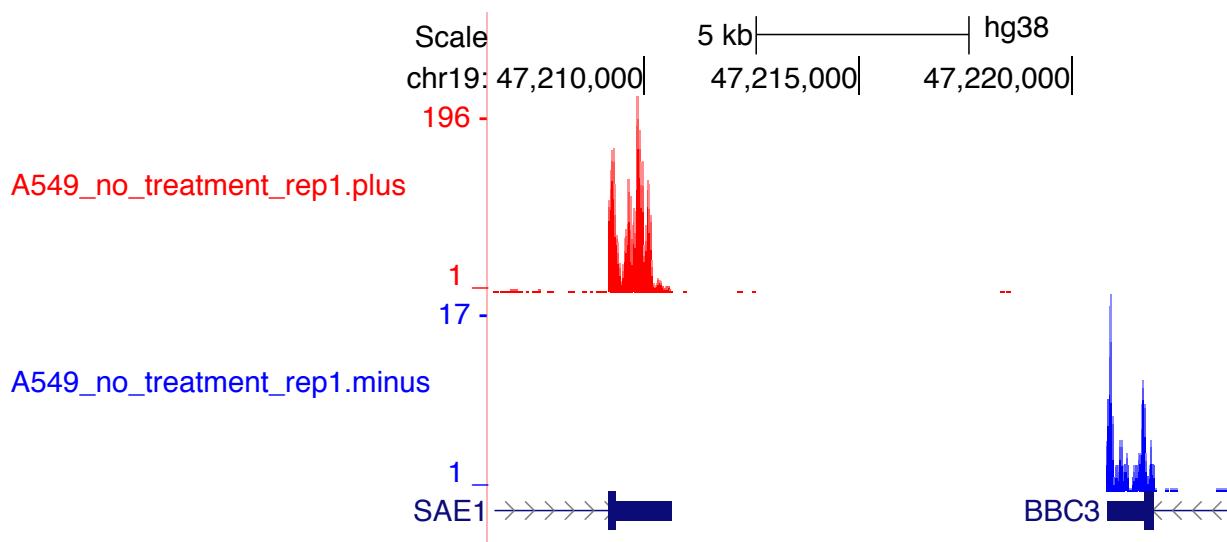


Figure 3: **Visualize your data in a browser.** Exploring your data in a browser can be the best way to troubleshoot and confirm or refute assumptions. For instance, looking at these tracks we can clearly see that the plus and minus tracks fall within the correct gene annotations. If we observed comparable densities on both the plus and minus exons, then it would be clear that the data was not strand-specific. If we saw that stranded tracks were inverted (i.e. the minus track overlaps the plus gene annotation), then we would have to change the `-strand` option in `genomeCoverageBed` in section 3.6. In fact, I could not find the method that ENCODE used to generate these libraries, so I had to empirically determine whether the reverse complement template or coding strand was sequenced.

4 Counting reads in genes

4.1 Installing HTSeq

```
pip3 install HTSeq
```

4.2 Count reads read in gene annotations

We use htseq-count to bin reads into genomic features (i.e. genes). The options below include -r which tells the program the bam file is sorted by position, -f which tells the program the input is a BAM file, and --stranded tells the program the strand information of the file. Notice the .gtf file is included. The output of this command is a .txt file that contains the gene counts for that sample. We will get one .txt file per sample. We confirmed that the data is stranded, but even if I know everything about the experiment, is not intuitive for me to know whether to specify yes or reverse for the --stranded argument.

To ensure that I use the correct option, I run both options and observe the results. Whatever option has more reads in the vast majority of the genes is correct. I appreciate that this is not clean, but these empirical checks are often necessary to ensure the data is properly processed. The biggest mistakes that I have observed in genomics could have been avoided if the bioinformatician looked at the output of analyses and confirmed that the result is expected or, at minimum, reasonable.

The output counts file has a footer (see the tail output) that we want to ignore in the counting. The grep -v command prints all the lines without __ to *stdout*, and we use awk to count the second column. I have been implementing awk code for ~15 years and I use *the Google* if I am not copying and pasting previously *Googled* commands. It's no *Ask Jeeves*, *Yahoo!*, or *Bing*, but I cannot overemphasize how useful *the Google* can be.

The loop is uploaded to GitHub: https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section4.2.sh.

Are there arguments in this loop that would be more appropriately initialized at the beginning of the script as opposed to hard coded?

```
name=A549_30min100nMdex_rep1.rmPCRdup.bam
htseq-count -r pos -f bam --stranded=yes $name gencode.v33.annotation.gtf > $name.gene.counts.yes.txt
htseq-count -r pos -f bam --stranded=reverse $name gencode.v33.annotation.gtf > $name.gene.counts.reverse.txt

tail $name.gene.counts.yes.txt
tail $name.gene.counts.reverse.txt

echo yes
grep -v __ $name.gene.counts.yes.txt | awk '{sum+=$2} END{print "sum=",sum}'
echo reverse
grep -v __ $name.gene.counts.reverse.txt | awk '{sum+=$2} END{print "sum=",sum}'

for i in *rmPCRdup.bam
do
    name=$(echo $i | awk -F".rmPCRdup." '{print $1}')
    echo $name
    htseq-count -r pos -f bam --stranded=yes $i gencode.v33.annotation.gtf > $name.gene.counts.txt
done
```

Use bash commands to take a look at the *.gene.counts.txt files. How many lines does each file contain? Are the first columns identical between files? Note the previously mentioned footer. What do the ENSG000* names represent and where are they found in the gtf annotation file?

5 Running jobs on a cluster with SLURM

After you retrieve the data, one can use the high performance computing cluster to streamline the routine analyses. We combined the alignment, generation of browser tracks, and counting in genes into a single bash script. I perform the following chunk in an interactive shell as opposed to submitting a job, but I make sure to keep the code for reference.

```
#ssh -Y mjt7y@rivanna.hpc.virginia.edu

module spider sratoolkit/2.9.1
module load sratoolkit/2.9.1

cd /scratch/mjt7y

fasterq-dump SRR5093037
fasterq-dump SRR5093036
fasterq-dump SRR5093038
fasterq-dump SRR5093039

fasterq-dump SRR5093268
fasterq-dump SRR5093265
fasterq-dump SRR5093267
fasterq-dump SRR5093266

mv SRR5093037_1.fastq A549_no_treatment_rep1_PE1.fastq
mv SRR5093036_1.fastq A549_no_treatment_rep2_PE1.fastq
mv SRR5093038_1.fastq A549_no_treatment_rep3_PE1.fastq
mv SRR5093039_1.fastq A549_no_treatment_rep4_PE1.fastq

mv SRR5093268_1.fastq A549_30min100nMdex_rep1_PE1.fastq
mv SRR5093265_1.fastq A549_30min100nMdex_rep2_PE1.fastq
mv SRR5093267_1.fastq A549_30min100nMdex_rep3_PE1.fastq
mv SRR5093266_1.fastq A549_30min100nMdex_rep4_PE1.fastq

mv SRR5093037_2.fastq A549_no_treatment_rep1_PE2.fastq
mv SRR5093036_2.fastq A549_no_treatment_rep2_PE2.fastq
mv SRR5093038_2.fastq A549_no_treatment_rep3_PE2.fastq
mv SRR5093039_2.fastq A549_no_treatment_rep4_PE2.fastq

mv SRR5093268_2.fastq A549_30min100nMdex_rep1_PE2.fastq
mv SRR5093265_2.fastq A549_30min100nMdex_rep2_PE2.fastq
mv SRR5093267_2.fastq A549_30min100nMdex_rep3_PE2.fastq
mv SRR5093266_2.fastq A549_30min100nMdex_rep4_PE2.fastq

gzip *fastq
```

5.1 BASH script

This script is called by the SLURM script below:

https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/200402_rna_seq.sh.

```
#!/bin/bash

idx=hg38
spliceFile=gencode.v33.splice.txt
processors=10

wget https://raw.githubusercontent.com/DaehwanKimLab/hisat2/master/hisat2_extract_splice_sites.py
wget ftp://ftp.ebi.ac.uk/pub/databases/gencode/Gencode_human/release_33/gencode.v33.annotation.gtf.gz
wget http://hgdownload.soe.ucsc.edu/goldenPath/hg38/bigZips/hg38.fa.gz
```

```

gunzip hg38.fa.gz
hisat2-build hg38.fa ${idx}

gunzip gencode.v33.annotation.gtf.gz
python3 hisat2_extract_splice_sites.py gencode.v33.annotation.gtf > ${spliceFile}

for i in *PE1.fastq.gz
do
    name=$(echo $i | awk -F"_PE." '{print $1}')
    echo $name
    hisat2 -p ${processors} -x ${idx} --rna-strandness FR --known-splicesite-infile ${spliceFile} \
        -1 ${i} -2 ${name}_PE2.fastq.gz | samtools view -bS - | samtools sort -n - -o $name.sorted.bam
    samtools fixmate -m $i - | samtools sort - | samtools markdup -rs - $name.rmPCRdup.bam
    htseq-count -r pos -f bam --stranded=yes $i gencode.v33.annotation.gtf > $name.gene.counts.txt
    samtools view -b -f 0x40 $i | samtools sort - -o ${name}.PE1.sorted
    samtools view -b -f 0x80 $i | samtools sort - -o ${name}.PE2.sorted
    genomeCoverageBed -bg -strand - -split -ibam ${name}.PE1.sorted | sortBed -i - > $name.pe1.minus.bedGraph
    genomeCoverageBed -bg -strand + -split -ibam ${name}.PE1.sorted | sortBed -i - > $name.pe1.plus.bedGraph
    genomeCoverageBed -bg -strand - -split -ibam ${name}.PE2.sorted | sortBed -i - > $name.pe2.minus.bedGraph
    genomeCoverageBed -bg -strand + -split -ibam ${name}.PE2.sorted | sortBed -i - > $name.pe2.plus.bedGraph
    cat $name.pe1.plus.bedGraph $name.pe2.plus.bedGraph | sortBed -i - > $name.merged.plus.bedGraph
    touch temp.txt
    echo "track type=bedGraph name=$name.plus color=255,0,0 visibility=2" >> temp.txt
    cat temp.txt $name.merged.plus.bedGraph > $name.plus.bedGraph
    rm temp.txt
    cat $name.pe1.minus.bedGraph $name.pe2.minus.bedGraph | sortBed -i - > $name.merged.minus.bedGraph
    touch temp.txt
    echo "track type=bedGraph name=$name.minus color=0,0,255 visibility=2" >> temp.txt
    cat temp.txt $name.merged.minus.bedGraph > $name.minus.bedGraph
    rm temp.txt
    rm *.merged.*us.bedGraph
    rm *.pe*.*us.bedGraph
    gzip $name.minus.bedGraph
    gzip $name.plus.bedGraph
done

```

5.2 SLURM script

This loads all the modules, specifies the job, and calls the processing script:

https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/200402_rna_slurm.slurm

```

#!/bin/bash
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=10
#SBATCH -t 72:00:00
#SBATCH --output=bioc8145_guertin_lecture
#SBATCH --partition=standard
#SBATCH -A bioc8145

module load gcc/7.1.0
module load hisat2/2.1.0
module load bedtools/2.26.0
module load samtools/1.10
module load htslib/1.10.2
module load bioconda/py3.6

./200402_rna_seq.sh

```

5.3 Job Submission

Change the file permission to executable using `chmod +x` and use the `sbatch` command to run the job. This analysis workflow is done in series one file at a time and it will take about 15 hours. Alternatively, one could write a script that generates a bash file for each `fastq` file pair and submit these all in parallel to different nodes on the cluster.

```
#make the shell script executable  
chmod +x 200402_rna_seq.sh
```

```
#run the job on the HPC cluster  
sbatch 200402_rna_slurm.slurm
```

6 Using R for analyses

We will use R to further analyze and plot the data. You must preinstall several libraries including *lattice*, *DESeq2*, *dplyr*, *fgsea*, *ggplot2* and *gage* (optional). I am not going into details for installing packages in R, but I am happy to try and help if you attempt to install these on your computer and fail. The first few lines of the chunk load in the libraries, and then I have you load a set of functions from my GitHub repository. I typically work locally in R. If I am performing memory-intensive operations I will use Rivanna for specific steps. Please note that I will not be able to help navigate RStudio, as I started using R in the console and I never adopted RStudio when it came out.

6.1 Installing R packages on Rivanna

Installing packages on a cluster is always more of a pain, but competent sysadmins, like we have for Rivanna, can be incredibly helpful. I submitted a help ticket and received assistance from Jacalyn Huband (jmh5ad) to install all the libraries on Rivanna. You should be able to use this script:

```
https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/install_pkgs.R
```

```
wget https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/install_pkgs.R
```

```
module load gcc openmpi R/3.6.1
```

```
R
```

6.2 Importing counts data into R

After defining the directory of interest, which contains my **.gene.counts.txt* files, I wrote a simple loop that will run through each **.gene.counts.txt* file and generate a data frame that has all the gene identifiers as row names, the experimental conditions as column names, and the corresponding counts at each column/row intersect. You can deconstruct the code to determine what is happening and I included comments. This *for loop* is something that you may want to call many times if you start analyzing a lot of RNA-seq data, so you can make this routine into a function.

```
https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section6.2.R
```

```
# if on rivanna
#module load gcc openmpi R/3.6.1
#you only need to run the next line once to install the packages
source('install_pkgs.R')

library(lattice)
library(DESeq2)
library(dplyr)
library(fgsea)
library(ggplot2)
library(gage)
library(limma)
source('https://raw.githubusercontent.com/mjg54/znf143_pro_seq_analysis/master/docs/ZNF143_functions.R')

path.dir='/scratch/mjg7y/'
file.suffix='.gene.counts.txt'
setwd(path.dir)

#keep track of the experiment names in this list so I can go back and label the columns
vec.names = c()
#we only want generate a data frame label the rows for the first file we loop through
count = 0
#here we will loop through our suffix-defined files in the defined directory
for (txt in Sys.glob(file.path(path.dir, paste0("*", file.suffix)))) {
```

```

count = count + 1
#this complicated code simply splits strings to extract the experiment name from the file
experiment.name = strsplit(strsplit(txt,
  "/")[[1]][length(strsplit(txt, "/")[[1]])], file.suffix)[[1]][1]
print(experiment.name)
#only do this for the first file, since count is greater than 1 with every other file
if (count == 1) {
#generate a data frame with the gene row names and no columns
  all.counts = data.frame(row.names = read.table(txt), ,1)
}
#add the experiment name to a growing list of column names
  vec.names = c(vec.names, experiment.name)
#for each file (including the first file) add a column with the counts information
  all.counts = cbind(all.counts, data.frame(read.table(txt), ,2))
}

#the last 5 lines are not gene counts, so delete them
all.counts = all.counts[1:(nrow(all.counts) - 5),]
#name the columns
colnames(all.counts) = vec.names

head(all.counts)

dim(all.counts)

```

6.2.1 Defining a function in R

The previous routine is formalized into a function. You can change the inputs to a function depending upon how you labeled the suffix of the counts files and where they are located on your computer.

https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section6.2.1.R

```

#remove the objects path.dir and file.suffix
rm(path.dir, file.suffix)

#these objects no longer exist
print(path.dir)
print(file.suffix)

#define a function
htseq.to.counts.df <- function(path.dir, file.suffix = '.gene.counts.txt') {
  vec.names = c()
  count = 0
  for (txt in Sys.glob(file.path(path.dir, paste0("*", file.suffix)))) {
    count = count + 1
    experiment.name = strsplit(strsplit(txt,
      "/")[[1]][length(strsplit(txt, "/")[[1]])], file.suffix)[[1]][1]
    print(experiment.name)
    if (count == 1) {
      all.counts = data.frame(row.names = read.table(txt), ,1)
    }
    vec.names = c(vec.names, experiment.name)
    all.counts = cbind(all.counts, data.frame(read.table(txt), ,2))
  }
  all.counts = all.counts[1:(nrow(all.counts) - 5),]
  colnames(all.counts) = vec.names
  return(all.counts)
}

```

```

all.counts.from.func = htseq.to.counts.df(path.dir = '/scratch/mjg7y/' ,
                                         file.suffix = '.gene.counts.txt')

#are the results identical?
identical(all.counts.from.func, all.counts)

```

6.3 Convert the Ensembl Gene (ENSG) name into a gene symbol

I prefer to have the gene names be more interpretable, so I convert the Ensembl Gene (ENSG) ID to the gene symbol.

https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section6.3.R

```

gencode.all = read.table('gencode.v33.annotation.gtf', sep='\t', header =F)

# the following few lines take some time to execute.
# this could be quicker with the use of data.table as opposed to data.frame
gencode.gene.names = data.frame(sapply(strsplit(sapply(strsplit(
  as.character(gencode.all[,9]), 'gene_name '), "[", 2), ";"), "[", 1)))

gencode.id.names = data.frame(sapply(strsplit(sapply(strsplit(
  as.character(gencode.all[,9]), 'gene_id '), "[", 2), ";"), "[", 1)))

gencode.key = cbind(gencode.gene.names, gencode.id.names)

gencode.key = gencode.key[!duplicated(gencode.key[,2]),]

rownames(gencode.key) = gencode.key[,2]

colnames(gencode.key) = c('gene', 'id')

#optional to save intermediate files
save(gencode.key, file = 'gencode.key.Rdata')

#why do I check the dimensions?
dim(all.counts)
dim(gencode.key)

merged.counts.pre = merge(all.counts, gencode.key, by="row.names", all.x=F)

merged.counts = merged.counts.pre[, c(2:(ncol(merged.counts.pre) - 2))]

rownames(merged.counts) = make.names(merged.counts.pre$gene, unique=TRUE)

#Optional exercise:
#Make this code chunk into a function.

```

6.4 Differential expression with DESeq2

DESeq (Anders and Huber, 2010; Love *et al.*, 2014) is an R library with the tools for differential gene expression analysis based on the negative binomial distribution (<http://www.bioconductor.org/packages/release/bioc/html/DESeq2.html>).

Refer to the DESeq2 vignette (<http://bioconductor.org/packages/release/bioc/vignettes/DESeq2/inst/doc/DESeq2.pdf>) and publication (Love *et al.*, 2014) for a detailed explanation of the processes and statistics.

The workflow below follows the DESeq2 vignette.

https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section6.4.R

```
#set the no treatment as the reference level (i.e. what happens upon treatment)
sample.conditions = factor(sapply(strsplit(as.character(colnames(merged.counts)),
  '_rep'), '[' , 1), levels=c("A549_no_treatment", "A549_30min100nMdex"))

#Note:
# Although not covered herein,
# I typically use these size factors to normalize bedGraph files.
# After I normalize each file (-scale in genomeCoverageBed),
# I convert to bigWig using bedGraphToBigWig
# and use bigWigMerge to make a signal representative track per condition.
# estimateSizeFactorsForMatrix(merged.counts)

deseq.counts.table = DESeqDataSetFromMatrix(merged.counts,
                                             as.data.frame(sample.conditions), ~ sample.conditions)

dds = DESeq(deseq.counts.table)

normalized.counts.rna = counts(dds, normalized=TRUE)

rld = rlog(dds, blind=TRUE)

pca.plot = plotPCA(rld, intgroup="sample.conditions", returnData=TRUE)
plotPCALattice(pca.plot, file = 'PCA_A549_dex_lattice.pdf')

#it is good to save the R session as you go.
#I name the image as a time stamp
save.image(file = '200404_R_mjg_1238pm.Rdata')
```

Principal components analysis (PCA) is a statistical method to reduce highly complex data into a limited number of dimensions that we can easily visualize. *A priori*, we expect that the samples separate based on conditions. The samples do not separate in this reasonable and expected manner. It is your responsibility to figure out the biological or technical feature that is driving the principal components. Here it is clear that replicate 4 is dominating PC1. We can hypothesize that something unique to replicate 4 is underlying PC1, which accounts for over half of the variance.




Figure 4: **PCA of RNA-seq analysis.** This shows that replicate 4 is dominating principal component 1, which is unexpected because the biology of interest is the dexamethasone treatment.

Visualize the DESeq2 dataframe with `results()`. I added a step to include a categorical column for different gene classes. The categories are defined by the multiple testing corrected p-value and and a log₂ fold change. How many genes are considered activated at the specified FDR (0.1)?

```
DE.results.all.4reps = results(dds)

head(DE.results.all.4reps)

DE.results.all.4reps.lattice =
  categorize.deseq.df(DE.results.all.4reps,
    fdr = 0.1, log2fold = 0.0, treat = 'Dexamethasone')

head(DE.results.all.4reps.lattice)

activated.all = DE.results.all.4reps.lattice[DE.results.all.4reps.lattice$response ==
  'Dexamethasone Activated',]

print(activated.all)
```

Take a look at the `DE.results.all.4reps.lattice` data frame and you will notice that no genes are differentially expressed. This can occur if the variance between replicates exceeds the variance between conditions. If we hypothesize that replicate 4 is driving the variation between replicates, then we would still expect that comparing the log₂ fold change and maybe the p-value to be meaningful. I would expect them to reflect real biology because the change when one adds dexamethasone may be consistent, even though the starting expression levels are clearly different for replicate 4. These differences in the basal expression will account for a lot of the variance and the differences will be common to both the untreated and Dexamethasone treated replicate 4, which would account for a high degree of variance within each experimental condition. The corrected p-values may not be significant, but the genes that change upon dexamethasone treatment should have a lower p-value.

6.5 Gene Set Enrichment Analysis

So we can use gene set enrichment analysis of a ranked gene set to determine what genes are enriched in our ranked set (Subramanian *et al.*, 2005). Recall that no genes are significantly activated or repressed, so we cannot use an overlap statistical analysis (e.g. Fisher's Exact). While categorically thresholding activated and repressed genes is useful, the threshold are biologically arbitrary and the quantitative nature of fold change is lost. GSEA employs rank information without the need for thresholding. Search around the Broad's GSEA site to learn more about the process. The following is a list of defined gene sets: <https://data.broadinstitute.org/gsea-msigdb/msigdb/release/>. Note the different collections to which you can compare in Figure 6 (<https://www.gsea-msigdb.org/gsea/msigdb/index.jsp>).

The MSigDB gene sets are divided into 8 major collections:

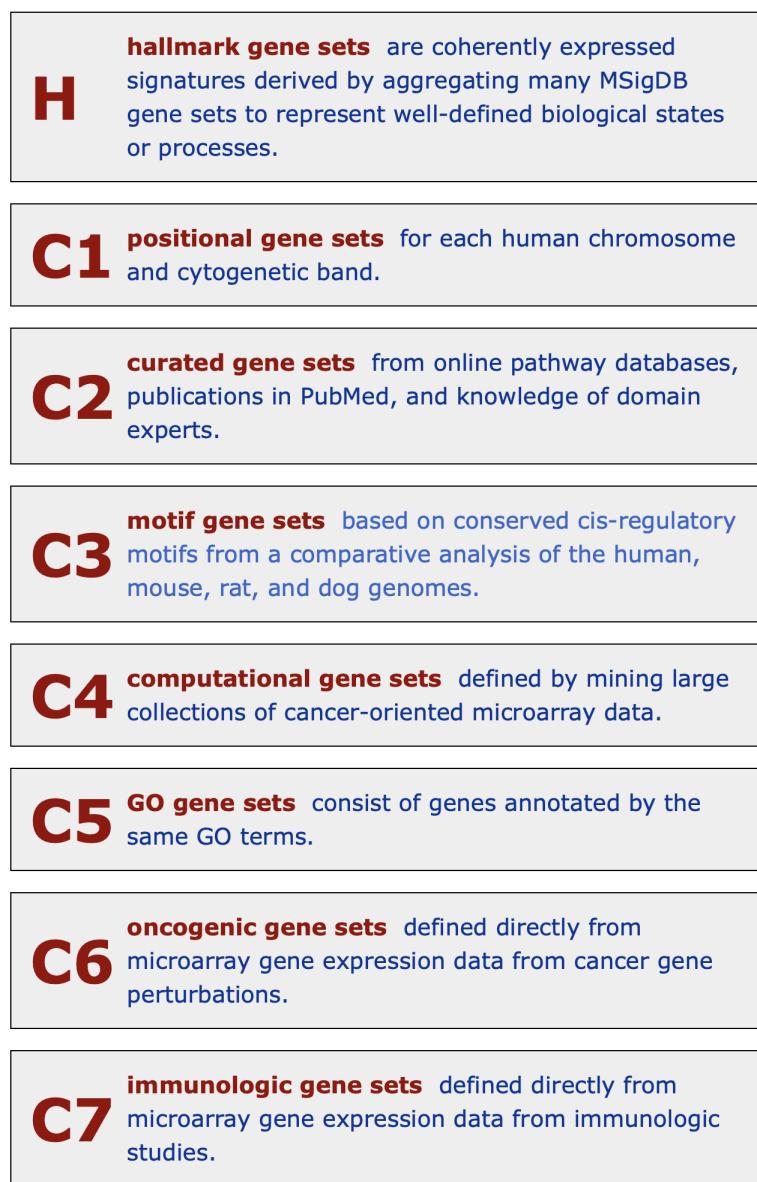


Figure 5: **Gene Set Enrichment Gene Collections.** These gene set enrichment collections are lifted directly from: <https://www.gsea-msigdb.org/gsea/msigdb/index.jsp>.

The first step is to order a gene list. One can justify ordering by p-value or fold change. Using either, I kept getting a warning: *There are ties in the preranked stats (XX% of the list). The order of those tied genes will be arbitrary, which may produce unexpected results.* The XX% was very high, so I multiplied the fold change by the $-\log_{10}$ p-value and used this to rank order the genes. You are encouraged to use other ranking methods to determine if any conclusions are robust to arbitrary ranking criteria. We will use fgsea in R for GSEA: <https://bioconductor.org/packages/release/bioc/html/fgsea.html>.

Analytical positive (and negative) controls can be as important as experimental controls. I use these as coherence check, especially when I am implementing some black box analysis tools that I do not fully understand. I manually typed table 2 (https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/G0_positive_ctrl.gmt) from a previous Dexamethasone-treated RNA-seq publication (Reddy *et al.*, 2009). The first author from this seminal molecular genomics paper is the PI who deposited the ENCODE data that we are currently analyzing.

https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section6.5.R

```
#you need to order the gene list:
gene.list= 2^DE.results.all.4reps.lattice$log2FoldChange *
           -log(DE.results.all.4reps.lattice$pvalue, base = 10)

names(gene.list) = row.names(DE.results.all.4reps.lattice)

#rank order
gene.list = sort(gene.list, decreasing = TRUE)

#remove NA entries
gene.list = gene.list[!is.na(gene.list)]

#remove duplicated genes
gene.list = gene.list[!duplicated(names(gene.list))]

head(gene.list)

# what file do you want to use?
#the more categories you have, the higher the multiple testing corrected test statistic

msigdb.url = 'https://data.broadinstitute.org/gsea-msigdb/msigdb/release/7.0/'

#GMT file options
#GO_file= paste0(msigdb.url, "h.all.v7.0.symbols.gmt")
#GO_file= paste0(msigdb.url, 'msigdb.v7.0.symbols.gmt')
#GO_file= paste0(msigdb.url, 'c5.all.v7.0.symbols.gmt')
#GO_file= paste0(msigdb.url, 'c2.all.v7.0.symbols.gmt')

#positive control
#check the content (second column) of the file to see where I manually copied it from
system('wget https://data.broadinstitute.org/gsea-msigdb/msigdb/release/7.0/h.all.v7.0.symbols.gmt')
system('wget https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/G0_positive_ctrl.gmt')
system('cat h.all.v7.0.symbols.gmt GO_positive_ctrl.gmt > h.all.v7.0.symbols.plusGRresp.gmt')

GO_file = 'h.all.v7.0.symbols.plusGRresp.gmt'
myGO = gmtPathways(GO_file)

fgRes = as.data.frame(fgsea(pathways = myGO,
                           stats = gene.list,
                           minSize=5,
                           maxSize=600,
```

```
    nperm=10000))

fgRes[fgRes$padj < 0.1,] [order(fgRes$NES[fgRes$padj < 0.1], decreasing =TRUE),]

plotEnrichment(myGO[["GO_DEX_ACTIVATED_GENES_WITH_GR_IN_A549"]],  
  gene.list) + labs(title="GR-bound Dex. Activated Genes in A549 (Reddy, et. al., 2009)")
```

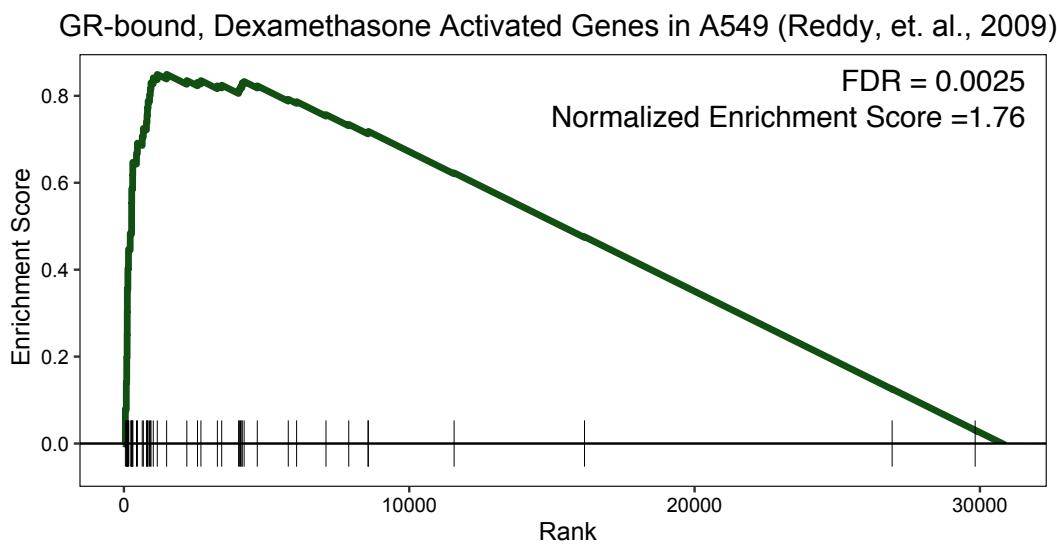


Figure 6: **Gene Set Enrichment plot.** The enrichment score is plotted against the genes ranked by the $-\log_{10}$ p-value multiplied by the fold change. Each gene in the class (in this case the 50 positive control genes) is represented by a black tick mark along the x-axis.

7 Unexpected results

I chose these data blindly, my criteria were stranded, paired-end, multiple replicates, publicly available, and two conditions separated by a short time interval. I was interested in these data because my group typically uses PRO-seq to measure nascent RNA and immediate changes in transcription. I was skeptical that we would see many genes changing expression and that the only genes that we would be able to detect activation would be short genes and those that were relatively lowly expressed during no-dex conditions. I hypothesized this because of the 30 minute treatment and because RNA-seq measures stable RNA. The two main barriers to look at changing expression are 1) the basal levels of RNA and 2) the gene length. Gene length is important because it takes a full minute for RNA polymerase to travel between 2-3 kilobases, so no gene that is 90 kilobases or more should be detectably changed at 30 minutes.

Now I realize that by not cherry-picking a polished and clean data set with straightforward results, I have the opportunity to teach you about the most challenging, time consuming, and rewarding component of bioinformatics: exploratory analyses to try and understand the unexpected. I believe that molecular biologists, or those with a deep understanding of biology, are at a distinct advantage in this realm.

I will take you step-by-step what I did to better understand these data and determine whether I believe dropping replicate 4 is justified.

7.1 Modeling batch effect of the replicates

Refer to the DESeq2 vignette for modeling batch effect: <http://bioconductor.org/packages/devel/bioc/vignettes/DESeq2/inst/doc/DESeq2.html>. Below we model the batch effect of each replicate, then we model two batches, operating under the assumption that an unknown variable in replicate 4 should be accounted for in the model. What *positive control* criteria can I use to evaluate whether the batch effect modeling is working? No treatment and Dexamethasone should separate by PCA and we should observe a few canonical dexamethasone activated genes as activated by GSEA. Since I am running the DESeq and GSEA code multiple times, I uploaded a very specific function that I will run below: https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/deseq.gsea.workflow.R.

The chunk below is found at: https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section7.1.R

```
github.week4 ='https://raw.githubusercontent.com/stefbekir/bioc8145/master/Guertin_week4/'

#This should be identical to the workflow that we performed outside the function:
#load function deseq.gsea.workflow
source(paste0(github.week4, 'deseq.gsea.workflow.R'))

#take a look at these vectors
sample.conditions = factor(sapply(strsplit(as.character(colnames(merged.counts)),
 '_rep'), '[', 1), levels=c("A549_no_treatment","A549_30min100nMdex"))

batch = factor(sapply(strsplit(as.character(colnames(merged.counts)),
 '_rep'), '[', 2))

gmt.file = 'h.all.v7.0.symbols.plusGRresp.gmt'

#recall merged.counts is the counts table

no.batch.all.four = deseq.gsea.workflow(merged.counts,
 batch = FALSE, sample.conditions,
 GO.file = gmt.file, name.analysis = 'No_batch_all_four_reps',
 treatment = 'Dexamethasone',
 fdr = 0.1, log2fold = 0.0)
```

```

head(no.batch.all.four[[1]])
head(no.batch.all.four[[2]])

#Apply a batch effect to each replicate

batch.all = factor(sapply(strsplit(as.character(colnames(merged.counts)),
 '_rep'), '[' , 2))

batch.all.four = deseq.gsea.workflow(merged.counts,
           batch = batch.all, sample.conditions,
           GO.file = gmt.file, name.analysis = 'Four_batches_all_four_reps',
           treatment = 'Dexamethasone',
           fdr = 0.1, log2fold = 0.0)

head(batch.all.four[[1]])
head(batch.all.four[[2]])

#is it reasonable to think that replicate 4 is a single batch and reps 1-3 are another?

batch.test = factor(c('batch1', 'batch1', 'batch1', 'batch2',
 'batch1', 'batch1', 'batch1', 'batch2'))

batch.all.two = deseq.gsea.workflow(merged.counts, batch = batch.test, sample.conditions,
           GO.file = gmt.file, name.analysis = 'Two_batches_all_four_reps',
           treatment = 'Dexamethasone', fdr = 0.1, log2fold = 0.0)

head(batch.all.two[[1]])
head(batch.all.two[[2]])

```

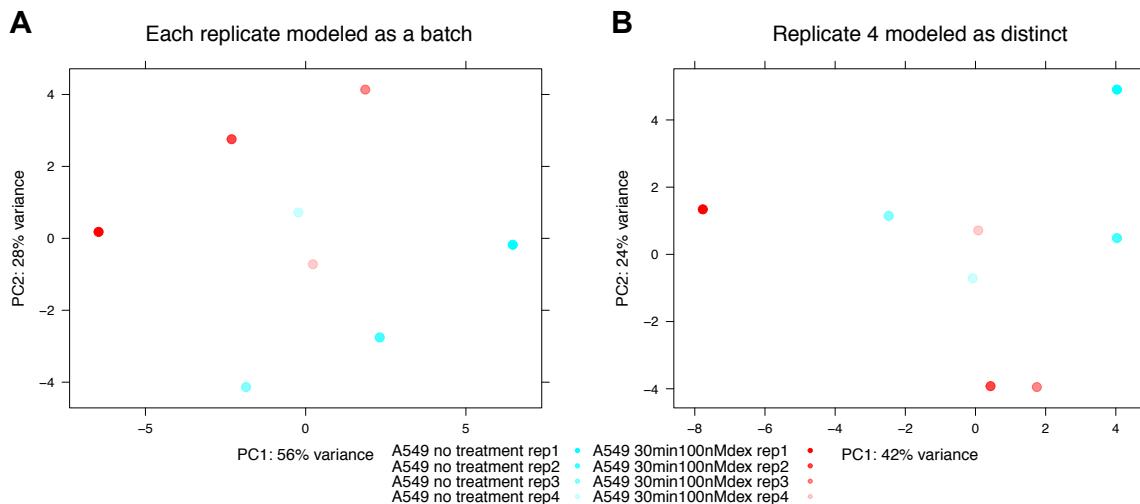


Figure 7: **PCA after batch effect modeling.** Incorporating batch effect residuals into the linear models employed by DESeq2 does not cleanly separate by experimental condition. Batch effect was modeled for each replicate (A) and as replicate 4 being a distinct batch compared to replicates 1-3 (B).

7.2 Exploratory analyses

In an attempt to get a better understanding of the data, I decided to perform a *leave one replicate out* as the comparison group analysis. I hypothesized that replicate 4 is distinct, so I started with comparing the other replicates to replicate 4.

https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section7.2.R

```
#we can change the sample conditions/comparison groups
sample.conditions.test = factor(c(rep("Replicates1-3", 3), "Replicate4",
                                 rep("Replicates1-3", 3), "Replicate4"),
                                 levels=c("Replicates1-3","Replicate4"))

deseq.counts.table = DESeqDataSetFromMatrix(merged.counts,
                                             as.data.frame(sample.conditions.test), ~ sample.conditions.test)

dds = DESeq(deseq.counts.table)

normalized.counts.rna = counts(dds, normalized=TRUE)

DE.results.rep4 = results(dds)

head(DE.results.rep4)

DE.results.rep4.lattice =
  categorize.deseq.df(DE.results.rep4,
    fdr = 0.1, log2fold = 0.0,
    treat = 'Replicate 4')

head(DE.results.rep4.lattice)

activated.rep4 = DE.results.rep4.lattice[DE.results.rep4.lattice$response ==
                                         'Replicate 4 Activated',]
dim(activated.rep4)
dim(DE.results.rep4.lattice[DE.results.rep4.lattice$response ==
                           'Replicate 4 Repressed',])

# Plot these results function
# I am in the minority in that I prefer lattice to ggplot

ma.plot.lattice <- function(ma.df, filename = 'file.name',
                            title.main = "Differential RNA-seq Expression",
                            col = c("grey90", "grey60", "red" , "blue"))
{
  pdf(paste("MA_plot_", filename, ".pdf", sep=''),
       useDingbats = FALSE, width=3.83, height=3.83);
  print(xyplot(ma.df$log2FoldChange ~ log(ma.df$baseMean, base=10),
               groups=ma.df$response,
               col= col,
               main=title.main, scales="free", aspect=1, pch=20, cex=0.5,
               ylab=expression("log"[2]^"RNA-seq change"),
               xlab=expression("log"[10]^"Mean of Normalized Counts"),
               par.settings=list(par.xlab.text=list(cex=1.1,font=2),
                                 par.ylab.text=list(cex=1.1,font=2))));
  dev.off()
}

ma.plot.lattice(DE.results.rep4.lattice, filename = 'A549_Replicate_4',
                title.main = "Differential Expression")
```

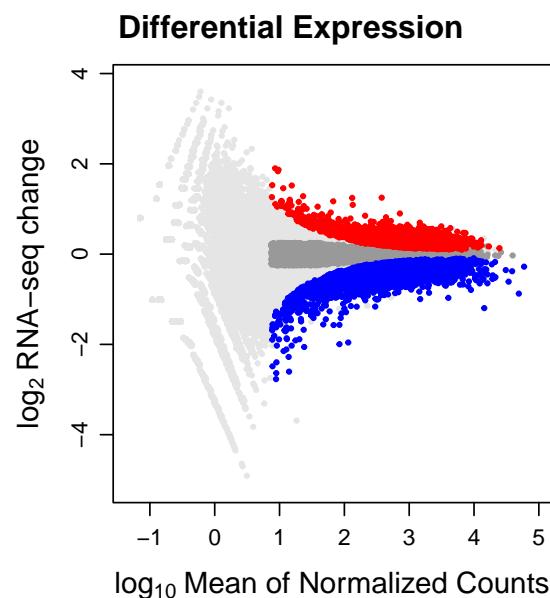


Figure 8: **Differential Expression of Replicate 4.** Each gene is a point that is colored by it's category. The dark grey points are confidently *unchanged*. The blue are *repressed* and red are *activated*. All other points are light grey. The red and blue genes are defined by the *fdr* and *log2fold* options with the *categorize.deseq.df* function. Comparing differential espression of Replicate 4 compared to the other three replicates indicates that about 3000 genes are both activated and repressed.

Comparing replicate 1-3 to replicate 4 resulted in 6019 genes differentially expressed (DE) at an FDR of 0.1, so I was curious if these number would be similar in other comparisons. 3 genes were DE when we used replicate 3 as the out group; 14 genes were DE when replicate 2 was the out group; 2 genes were DE when replicate 1 was the out group.

Taking a look at the gene set enrichment analysis when replicate 4 was considered the treatment group, the categories *ribosome*, *translation*, and *serum response* are found multiple times among the top pathways. Ribosome biogenesis is linked to nutrient availability (Thomas, 2000) and cells in culture react strongly to serum (Mahat and Lis, 2017). I hypothesize that the experimental conditions of these cells were distinct compared to the other replicates in the time since they were split or the lot number (and thus nutrient content) of the FBS that was used.

I do not know of a way to test this hypothesis, but this scenario highlights the importance of keeping a well annotated lab notebook that you can refer back to. Regardless, I think there is sufficient justification for excluding replicate 4 from the analyses. However, this is subjective and I would urge transparency in methods if you have to make these calls in a publication. One can use GitHub to track the step-by-step analyses that go into a manuscript and provide excruciating detail for these type of exploratory analyses that inform on good faith judgement calls that the researcher has to take.

7.3 Final analyses

Comparing dexamethasone treatment in replicates 1-3 indicates that canonical NF κ B signaling genes are changing expression and there are 31 activated and 11 repressed genes at an FDR of 0.1. Much is already known about the antagonism of glucocorticoid signaling and NF κ B signaling, so this result is consistent with known biology. https://github.com/stefbekir/bioc8145/blob/master/Guertin_week4/section7.3.R

```
sample.conditions.1to3 = factor(sapply(strsplit(as.character(colnames(merged.counts)),  
    '_rep'), '[' , 1)[c(1,2,3,5,6,7)],  
    levels=c("A549_no_treatment","A549_30min100nMdex"))  
  
merged.counts.1to3 = merged.counts[,c(1,2,3,5,6,7)]  
  
deseq.gsea.1to3 = deseq.gsea.workflow(merged.counts.1to3,  
    batch = FALSE, sample.conditions.1to3,  
    GO.file = 'h.all.v7.0.symbols.plusGRresp.gmt',  
    name.analysis = 'Dexamethasone',  
    treatment = 'Dexamethasone',  
    fdr = 0.1, log2fold = 0.0)  
  
ma.plot.lattice(deseq.gsea.1to3[[1]], filename = 'A549_RNA_dex',  
    title.main = "Differential Expression")
```

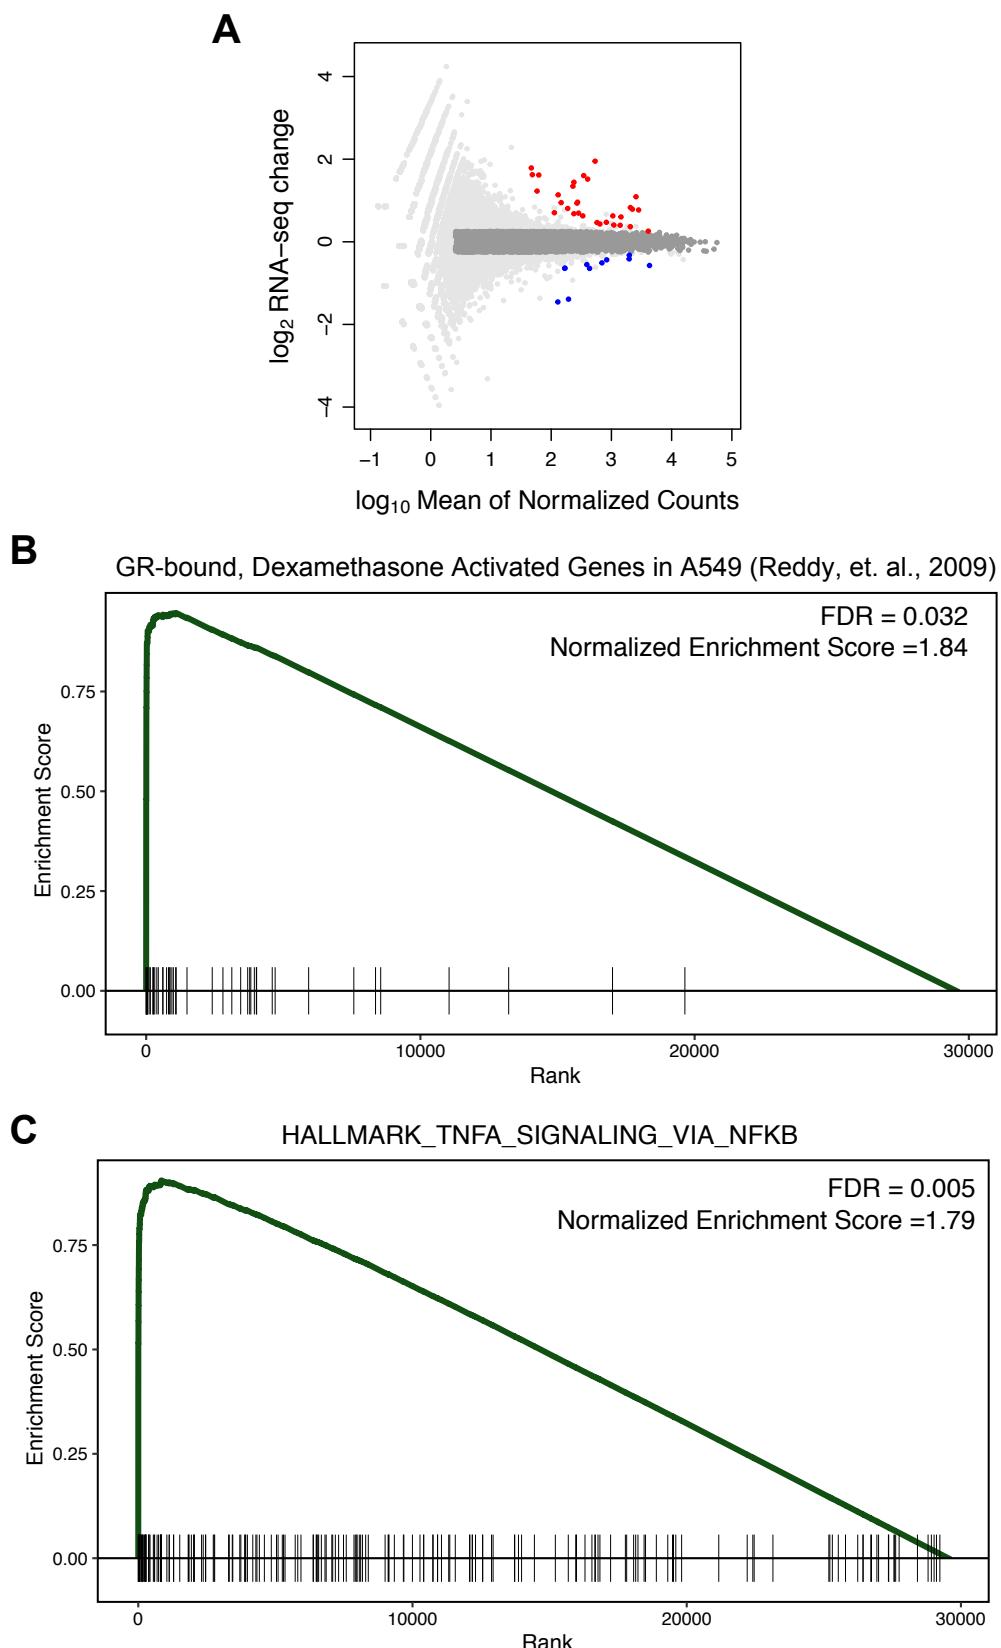


Figure 9: **MA plot and GSEA of final analysis.** A) An MA plot comparing no treatment and dexamethasone treated cells with replicate 4 excluded. B) The positive control gene set contains the highest normalized enrichment score. C) NF κ B signaling is the most enriched gene set outside the positive control.

References

- Anders S, Huber W (2010). "Differential expression analysis for sequence count data." *Genome biol*, **11**(10), R106.
- Karolchik D, Barber GP, Casper J, Clawson H, Cline MS, Diekhans M, Dreszer TR, Fujita PA, Guruvadoo L, Haeussler M, et al. (2014). "The UCSC genome browser database: 2014 update." *Nucleic acids research*, **42**(D1), D764–D770.
- Kim D, Langmead B, Salzberg SL (2015). "HISAT: a fast spliced aligner with low memory requirements." *Nature methods*, **12**(4), 357–360.
- Kodama Y, Shumway M, Leinonen R (2011). "The Sequence Read Archive: explosive growth of sequencing data." *Nucleic acids research*, **40**(D1), D54–D56.
- Love MI, Huber W, Anders S (2014). "Moderated estimation of fold change and dispersion for RNA-Seq data with DESeq2." *bioRxiv*.
- Mahat DB, Lis JT (2017). "Use of conditioned media is critical for studies of regulation in response to rapid heat shock." *Cell Stress and Chaperones*, **22**(1), 155–162.
- Mi H, Muruganujan A, Ebert D, Huang X, Thomas PD (2019). "PANTHER version 14: more genomes, a new PANTHER GO-slim and improvements in enrichment analysis tools." *Nucleic acids research*, **47**(D1), D419–D426.
- Reddy TE, Pauli F, Sprouse RO, Neff NF, Newberry KM, Garabedian MJ, Myers RM (2009). "Genomic determination of the glucocorticoid response reveals unexpected mechanisms of gene regulation." *Genome research*.
- Sathyan KM, McKenna BD, Anderson WD, Duarte FM, Core L, Guertin MJ (2019). "An improved auxin-inducible degron system preserves native protein levels and enables rapid and specific protein depletion." *Genes & development*, **33**(19-20), 1441–1455.
- Subramanian A, Tamayo P, Mootha VK, Mukherjee S, Ebert BL, Gillette MA, Paulovich A, Pomeroy SL, Golub TR, Lander ES, et al. (2005). "Gene set enrichment analysis: a knowledge-based approach for interpreting genome-wide expression profiles." *Proceedings of the National Academy of Sciences*, **102**(43), 15545–15550.
- Thomas G (2000). "An encore for ribosome biogenesis in the control of cell proliferation." *Nature cell biology*, **2**(5), E71–E72.
- Wang Z, Civelek M, Miller C, Sheffield N, Guertin M, Zang C (2018). "BART: a transcription factor prediction tool with query gene sets or epigenomic profiles." *Bioinformatics (Oxford, England)*, **34**(16), 2867.